

Autoencoder reconstruction accuracy by region (test set, overall AAR = 97.73%)
Frameworks reconstruct near-perfectly; the CDRs (shaded) are hardest — the variable region the latent must compress

